

A Kalman filter model of decision making

One interesting aspect of decision making in humans and other animals, is that it appears to place a value in information. This is often described using the example of someone living in a village with two night clubs. Imagine this person goes out one Friday and chooses one of the two night clubs, The Oasis, say and has a terrible night, awful music, fighting on the dance floor and so on. The next Friday night they go to the other night club, The Warwick say and have a great time, everyone dancing at the disco, having a great time. In an unchanging world, the best thing to do from then on would be to go to The Warwick, it is the better night club after all.

However, if they did this they might loose out, The Warwick might go down hill, employ a bad DJ and aggressive bouncers and stop replacing the little mirrors in their disco ball when they fall out. At the same time, The Oasis might get its act together, might stop playing Europop and discourage dance-floor fighting. The village person at this point would be going to the wrong night club. Obviously the correct thing to do is to sometimes make a decision that is not based on the optimal predicted reward.

There is now a very general theory of decision making call the Free Energy Principle that attempts to model the value of knowledge. Here, we will look at a more restricted attempt, in Daw et al. (2006) to model what is called *exploration versus exploitation*. In this paper they examine a number of different models¹, but the one that seems to work best is based on a so-called *soft-max* rule.

If there are n alternative actions, with predicted reward r_i for the i th action, then a soft-max rule says that the i th action is selected with probability

$$p_i = \frac{e^{\beta r_i}}{\sum_j e^{\beta r_j}} \quad (1)$$

Here β is a parameter which determines the amount of exploration, for large β then the action with the largest reward will have a p_i near to one, whereas for $\beta = 0$ all the probabilities are equal; this β should respond to the predictability of the environment with β increasing if the world is more predictable. β might be controlled by neuromodulation. However, that is another question, the one of interest here is to fit this model and to compare it to other models,

¹Note that most of the actual details for this paper appear in the supplementary materials section

such as one where the probabilities are

$$p_i = \begin{cases} 1 - \epsilon & r_i \text{ is the biggest } r_j \text{ value} \\ \epsilon/n & \text{otherwise} \end{cases} \quad (2)$$

The difficulty is to work out what the predicted rewards, r_i , are.

An n -armed bandit



Figure 1: A one-armed bandit. [Image from ebay]

Daw et al. (2006) describes a so-called n -armed bandit task, named after the ‘one-armed bandit’ gambling machine as in Fig. 2. In the experiment a subject has a choice of four buttons to press. Each button has a reward, but the amount is different for each button and the amount for each button varies with time. In fact the mean reward is changed from trial to trial according to a discrete Ornstein-Uhlenbeck process, that is a kind of mean-reverting Gaußian random walk:

$$\mu_i(t+1) = \lambda\mu_i(t) + (1-\lambda)\mu_0 + v \quad (3)$$

where t is the trial number, μ_0 is the mean value for the mean, λ is a parameter controlling how rapidly the reward heads towards \bar{r} and v is drawn independently for each trial from a zero mean Gaußian with variance σ_v^2 . The actual reward for a trial is then chosen to be

$$r_i(t+1) = \mu_i(t+1) + w \quad (4)$$

where w is also a zero mean Gaussian with variance σ_w^2 . A schematic for the experiment is showing in Fig. ??.

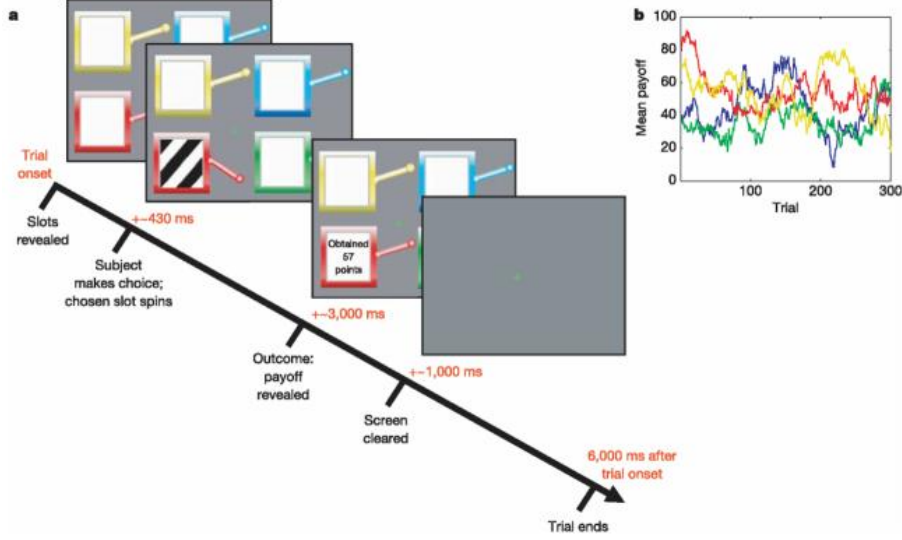


Figure 2: A schematic for the n -armed bandit task, **b** shows examples of how the mean rewards $\mu_i(t)$ evolve. [Image from (Daw et al., 2006)]

Now the idea presented in the paper is that the subject estimates μ_i using a one-dimensional Kalman filter. Thus, if $\tilde{\mu}_i(t)$ is the predicted reward for button i at trial t . According to their model the subject has an estimates of λ and μ_0 which we will call $\tilde{\lambda}$ and $\tilde{\mu}_0$. Thus from one trial to the next the estimate $\tilde{\mu}_i$ is updated according to

$$\tilde{\mu}_i^-(t+1) = \tilde{\lambda}\mu_i(t) + (1 - \tilde{\lambda})\tilde{\mu}_0 \quad (5)$$

and the estimate of the uncertainty of this estimate is updated as

$$\tilde{\sigma}_i^{2-}(t+1) = \tilde{\lambda}^2\tilde{\sigma}_i^2(t+1) + \sigma_v^2 \quad (6)$$

where in both equations the subscripted $-$ indicates this is the predicted value before taking into account the actual observation, the actual reward received from whichever button is pressed. For one of the buttons there is also an update based on the actual reward:

$$\tilde{\mu}_i(t+1) = \mu_i^-(t+1) + \kappa(t+1)[r(t+1) - \mu_i^-(t+1)] \quad (7)$$

and

$$\tilde{\sigma}_i^2(t+1) = [1 - \kappa(t+1)]\tilde{\sigma}_i^{2-}(t) \quad (8)$$

where $\kappa(t+1)$ is the Kalman gain

$$\kappa(t+1) = \frac{\tilde{\sigma}_i^{2-}(t+1)}{\sigma_w^2 + \tilde{\sigma}_i^{2-}(t+1)} \quad (9)$$

However, for the other three buttons there is no observation and so

$$\begin{aligned} \tilde{\mu}_i(t+1) &= \tilde{\mu}_i^-(t+1) \\ \tilde{\sigma}_i^2(t+1) &= \tilde{\sigma}_i^{2-}(t+1) \end{aligned} \quad (10)$$

In Daw et al. (2006) the Kalman filter model is fitted to the actual decision-making data.² This gives a value for the predicted average reward.

Exploration versus exploitation and the brain

Fitting the Kalman filter as two benefits, firstly it allows the authors to test the different models of decision making; as mentioned above, the one that works best is soft-max. It also allows the authors to distinguish between trials where the subjects appear to pick the button with the highest predicted reward and the trials where the subjects appear to explore, by picking a button different from that one. Figure 3 shows the brain region which activates distinguishes most strongly between exploitation and exploration trials: the frontopolar cortex seems to have higher BOLD signal for exploration trials compared to exploitation trials.

References

Daw, N. D., O'Doherty, J. P., Dayan, P., Seymour, B., and Dolan, R. J. (2006). Cortical substrates for exploratory decisions in humans. *Nature*, 441(7095):876.

²The fit is actually remarkably poor, it requires values for $\tilde{\lambda}$, for example, that are very different from λ . We will ignore this problem here, as the authors did in their paper; it seems not to prevent the model producing interesting results.




Figure 3: In **a** shows a significance map for the difference in BOLD activation between exploration and exploitation trials. For the yellow regions the difference in BOLD activation is different from zero with significance $p < 0.001$, for the red regions, with $p < 0.01$. **b** shows the average time course of the BOLD activation for exploration (red, mostly upper) and exploitation (blue, mostly lower) trials; roughly speaking, the actual data is marked by the black points, the lines are fitted to these points. [Image from (Daw et al., 2006)]