Neural Information Processing 2018/2019

University of BRISTOL

Brainbow (Litchman Lab)

Lecture 13

Neural circuits and learning:

Unsupervised learning:

Sparse coding and autoencoders

Extra lab session

This Wednesday (12-2pm, 5th of December)!

Sign up for help with course work:
https://doodle.com/poll/tbky2r8st4xxc3kb

Previously on *Neural Information Processing*...

Supervised learning and backprop:

Relies on a teaching signal

But, most sensory data is unsupervised

Every few milliseconds we experience a new input!

https://youtu.be/Plx9IJ1KOsg

But, most data animals experience is unsupervised

Unsupervised Learning:

Extracts useful representations of input

Outline

1. **Sparse coding**
2. **Sparse coding in the brain**
3. **Autoencoders**

Sparse coding

Sparse coding

- **Sparse coding** provides a framework for learning sparse representations of the input, typically through unsupervised learning.

- *Definition*: '**Sparse coding** is the representation of items by the strong activation of a relatively small set of neurons. For each stimulus, this is a different subset of all available neurons.'

Foldiak and Dominik, Scholarpedia (2008)

Sparse coding

Key properties of <u>sparse coding</u>:

- **Overcomplete dictionary *W*: d>m** (i.e. more dictionary elements than input elements)
- **Sparseness**: # of zero elements in ***v*** >> # non-zero elements in ***v***
- Overcomplete and sparse representation leads to **<u>richer features</u>**

Notes:

$$\|X\|_2^2 = \sum_{i=1}^{n} |x_i|^2$$

$$\|X\|_1 = \sum_{i=1}^{n} |x_i|$$

Sparse coding
L1-norm

$$\text{cost} = [\text{preserve information}] + \lambda[\text{sparseness}]$$

reconstruction cost

enforces sparseness

Sparse coding with L1-norm:

$$\text{argmin}_{W,V} = \|U - WV\|_2^2 + \lambda\|V\|_1$$

$$\text{subject to } \|W_i\|_2^2 \leq 1 \;\; \forall_i = 1,\ldots,m$$

avoids arbitrarily high ***W***

- To find a good sparse code we need to optimise both the dictionary, ***W*** and the representation ***V***
- This optimisation problem is convex in ***V*** and ***W*** separately, but not both. Solution: interleaved optimisation of ***V*** and ***W***
- Several methods exist to optimise **V** and **W** (we will discuss one)
- λ controls the tradeoff between sparseness and reconstruction

Sparse coding

Input: Natural images

Simple cell like features:

Sparse coding

Input: MNIST dataset (handwritten digits)

Sparse coding
Applications: denoising

Original

Sparse coding method

Sparse coding
Applications: image restoration

Original

Sparse coding method

Fig. 8. Underwater. Left: Image containing many high contrast edges, with text covering 16.19% of its area. Right: restored

Richard et al. VIIP 2001

Sparse coding
Applications: digital zooming

Original

Sparse coding method

Couzinie-Devy, 2010

Sparse coding in the visual system

Barlow (1981)

Sparse coding in the brain

Olshausen and Field 1996

New cost function:

new sparseness cost

$$\text{argmin}_{W,V} = \overbrace{\|U - WV\|_2^2 + \lambda \sum_i S(v_i)}^{\text{Cost}}$$

$$S(v_i) = \log(1 + v_i^2)$$

Notation in Olshausen and Field paper:

$$\text{argmin}_{\Phi,A} = \|I - \Phi A\|_2^2 + \lambda \sum_i C(a_i)$$

Olshausen and Field, Nature (1996)

Sparse coding in the brain

Olshausen and Field 1996

Compute *V* via gradient descent, and get dynamics for *v*:

$$\dot{v}_i = -\frac{\partial \text{Cost}}{\partial v_i}$$

$$= w_i U - \sum_{j \neq i} G_{ij} v_j - f_\lambda(v_i)$$

where

$$G_{ij} = w_i w_j$$

$$f_\lambda(v_i) = v_i + \lambda S'(v_i)$$

Olshausen and Field, Nature (1996)

Sparse coding in the brain

Olshausen and Field 1996

Potential network implementation of:

$$\dot{v}_i = w_i U - \sum_{j \neq i} G_{ij} v_j - f_\lambda(v_i)$$

Olshausen and Field, Nature (1996)

Sparse coding in the brain

Olshausen and Field 1996

With v we obtain the estimate, on a second phase we need to optimise the weights as:

$$\Delta w_i = -\eta \frac{\partial \text{Cost}}{\partial w_i}$$

reconstructed image

$$= \eta < [U - \hat{U}] v_i >$$

$$= \eta < [U - \sum_i^m w_i v_i] v_i >$$

Learning rule is local, and hebbian:

$$\Delta w \propto \underbrace{UV}_{\text{Hebbian term}} - \underbrace{WV^2}_{\text{post. alone}}$$

Note: Hebbian learning rules depend on presynaptic and postsynaptic activity, e.g.:

$$\Delta w = UV$$

Hebbian plasticity: Cells that fire together, wire together

Olshausen and Field, Nature (1996)

Sparse coding in the brain

Olshausen and Field 1996

Learned dictionaries/weights (simple cell-like)**:**

Olshausen and Field, Nature (1996)

Group discussion
groups of 2-3 (5min)

Can you think of advantages of using **sparse coding** in the brain?

Group discussion
groups of 2-3 (5min)

Can you think of other advantages of using **sparse coding** in the brain?

- **Good representations** that can be used for decision making/ discrimination
- Sparse coding is **energy efficient** (only a few neurons active at a time)
- **More robust** to perturbations than purely localist codes

Sparse coding

Through Hebbian learning

inhibitory synapses, *w*

$$\Delta w_{ij} = -\underline{\alpha(y_i y_j} - p^2)$$

Anti-Hebbian term
(because of the minus sign)

excitatory synapses, *q*

$$\Delta q_{ij} = \beta \underline{y_i(x_j} - q_{ij})$$

Hebbian term

Hebbian plasticity: Cells that fire together, wire together

Foldiak, Biological Cybernetics (1990)

Sparse coding

Through Hebbian learning

Input defined as random patterns of vertical and horizontal lines:

Learned features:

Foldiak, Biological Cybernetics (1990)

Sparse coding

Spiking neural network: SAILnet

SAILnet:

A spiking version of Foldiak 1990 network with inhibitory and excitatory synapses and local learning rules:

Zylberberg et al. PLoS CompBIo (2011)

Sparse coding
SAILnet

Captures the diversity of receptive fields observed in animals:

Zylberberg et al. PLoS CompBIo (2011)

Sparse coding across the animal kingdom:
from mammals to insets..

Sparse coding in fruit flies

Smell (odor) system in *Drosophila*

Masse et al. Current Biology (2009)

Sparse coding in Kenyon cells

Lin et al. Nature Neuroscience (2014)

Autoencoders

curiously.com

$$\text{argmin}_W = \|V - U\|_2^2$$

- **Autoencoders** also aim at reconstructing the input, but without an explicit constrain on sparsity (instead uses other forms of regularization)
- **Autoencoders** rely on a specific model *f* (i.e. neural network with bottleneck)
- If an **autoencoders** uses a sparsity constrain, then are equiv. to **sparse coding**

Autoencoders

Cat neuron

Autoencoder trained on **10 million images** from **YouTube** movies.

Showing that it is possible to learn high level feature detectors from unlabelled data (akin to grandmother cells).

Le et al. ICML (2012)

Summary

1. **Sparse coding** is an **optimisation** method that develops useful representations of the input
2. Olshausen and Field introduced a sparse coding framework that can be mapped onto the brain
3. **Sparse coding** develops **representations akin** to the brain
4. Autoencoders can also be used to extract useful representations, but do not (typically) rely on a sparsity cost and assume a specific neural network

References

Text books:

Natural Image Statistics, by Hyvarinen et al. 2009 [good reference on unsupervised learning from natural images]

Theoretical neuroscience: Dayan and Abbott 2001 [Unsupervised learning: chapter 3]

Deep Learning by Courville et al. 2015

Relevant papers:

- Foldiak, Biological Cybernetics (1990) [classical work on developing sparse representations with local Hebbian learning rules]
- Olshausen and Field, Nature (1996) [seminal work on sparse coding using natural images]
- Zylberberg et al., PLoS Computational Biology (2001) [more recent work on sparse coding in spiking networks with local learning rules]

Upcoming lectures

- **L10: Neural circuits and learning: introduction**
  - Visual processing
    - **L11: Visual cortex**
    - **L11: Convolutional neural networks**
  - Learning in the brain
    - L12: Supervised learning: The backpropagation algorithm/cerebellum
    - L13: Unsupervised learning: Sparse coding and autoencoders
    - **L14: Reinforcement learning: TD learning, Q learning, deep RL and dopamine**
  - Temporal processing in the brain
    - L14: Auditory cortex and recurrent neural networks
    - L15: Gated recurrent neural networks